

Week 3 Lab Session

Adam Wu, Wonje Yun

April 5, 2024

University of Chicago

1. Review
2. Assignment Submission Example
3. Review of Numba/MPI
4. GPU Jobs
5. Assignment Walk-through

Review of Access

- Github setup
- Midway code editing
- Set up the necessary libraries at the base environment

Assignment Submission

Assignment Submission Example

- `git clone` assignment repository to Midway3
- Create folders so that assignments are clear
- Create separate python files if necessary (do not have to attempt to answer all in one file)
- Write down your code, run the code and get results
- Write down the README as described in Canvas < Assignment 1 and link

Review of Numba/MPI

Review of Numba/MPI/sbatch

- In pyCC, function signs are crucial. Need to figure out your input and output type.
- What happens when `mpirun` a job and between `comm.Scatter` and `comm.Gather`?
 - Unless `rank` is specified, every node receives the input and returns output
 - To have only one output, put the `print` or `return` under

```
1 if rank == 0:  
2     # YOUR CODE HERE
```

- `.sbatch` file contains the configuration of HPC hardware the user is requesting, and the command that it is requested to run.

For more SLURM commands see also: [Slurm documentation](#)

MPI for Python - 1

- `MPI.COMM_WORLD`: set up communication for mpi
- `comm.Get_size()`: return the number of processes
- `comm.Get_rank()`: return the rank of this process
- `comm.Bcast()`: broadcast a message from one process to all other processes in a group
- `comm.Scatter()`: scatter data from one process to all other processes in a group
- `comm.Gather()`: gather together values from a group of processes

For more explanation: [mpi4py.MPI documentation](#) and [bcast/scatter/gather example](#)

- `comm.Bcast(buf, root=0)`: broadcast the data(`buf`) at rank 0 to all workers
- `comm.Scatter(sendbuf, recvbuf, root=0)`: scatter data(`sendbuf`) at rank 0 to `recvbuf` at each workers.
- `comm.Gather(sendbuf, recvbuf, root=0)`: gather the data(`sendbuf`) in each workers to the `recvbuf` at rank 0.
- If you want to do some global-like operation before scattering or after gathering, the operation should go under "`if rank == 0:`" statement.

Using GPU with OpenCL

- `pyopencl` gives a pythonic access to OpenCL API.
- As in `mpi4py`, we need codes within the python function to make it run as an `opencl` job.
- `cl.create_some_context()`: create a `opencl` context 'somehow' (a device is chosen in an implementation-defined manner).
 - The `Context` is a construct that manages the set of devices, memory and command-queues.
- `cl.CommandQueue()`: create a command queue using the specified context

```
1 import pyopencl as cl
2 ctx = cl.create_some_context()
3 queue = cl.CommandQueue(ctx)
```

- `pyopencl.array.Array`: The `numpy.ndarray` work-alike, so that it functions like an numpy array but operates on OpenCL.
[pyopencl.array.Array](#)
- `pyopencl.elementwise.ElementwiseKernel`: Performs element-wise computation on the `pyopencl.array.Array` object.
 - `context`: the OpenCL context
 - `arguments`: the variable argument of the operation (need to assign data types, for arrays add * before the variable)
 - `operation`: the element-wise operation to perform
 - `name`: the function name of the compiled kernel
 - `Ref`: [ElementwiseKernel](#)

ElementwiseKernel Code Example

```
1 import pyopencl.array as cl_array
2 from pyopencl.elementwise import ElementwiseKernel
3
4 ctx = cl.create_some_context()
5 queue = cl.CommandQueue(ctx)
6
7 a_g = cl.array.to_device(queue, a_np)
8 b_g = cl.array.to_device(queue, b_np)
9
10 lin_comb = ElementwiseKernel(ctx,
11     "float k1, float *a_g, float k2, float *b_g, float *res_g",
12     "res_g[i] = k1 * a_g[i] + k2 * b_g[i]",
13     "lin_comb")
14
15 res_g = cl.array.empty_like(a_g)
16 lin_comb(2, a_g, 3, b_g, res_g)
```

- `pyopencl.reduction.ReductionKernel`: Performs reduction operation (eg. count, sum, etc.) on the `pyopencl.array.Array` object.
 - `context`: the OpenCL context
 - `dtypeout`: the dtype of the output
 - `neutral`: initial value
 - `map_expr`: the entry operation of the vector
 - `reduce_expr`: the operation to perform on the result of `map_expr`
 - `arguments`: the input (needs at least one argument to be vector)
 - Ref: [ReductionKernel](#)

ReductionKernel Code Example

```
1 import pyopencl.array as cl_array
2 from pyopencl.reduction import ReductionKernel
3
4 ctx = cl.create_some_context()
5 queue = cl.CommandQueue(ctx)
6
7 a = cl_array.arange(queue, 400, dtype=numpy.float32)
8 b = cl_array.arange(queue, 400, dtype=numpy.float32)
9
10 krnl = ReductionKernel(ctx, numpy.float32, neutral="0",
11                        reduce_expr="a+b", map_expr="x[i]*y[i]",
12                        arguments="__global float *x, __global float *y")
13
14 my_dot_prod = krnl(a, b).get()
```

GPU sbatch Commands

For using GPUs, different sbatch commands are needed.

- **-time**: set a limit on the total run time of the job allocation
- **-nodes**: number of GPU nodes to use
- **-partition**: set gpu for Midway3
- **-gres**: the number of gpus to request (gpu:1 for 1 GPU usage)
- **-ntasks-per-node**: number of cores to drive each gpu
- **-mem-per-cpu**: minimum memory required per usable allocated CPU

For more, see also: [RCC user guide for SLURM](#) and [Slurm documentation](#)

GPU sbatch

```
1 #!/bin/bash
2
3 #SBATCH --job-name=gpu_mpi
4 #SBATCH --output=gpu_mpi.out
5 #SBATCH --error=gpu_mpi.err
6 #SBATCH --nodes=1 # using 1 GPU node
7 #SBATCH --ntasks-per-node=1 # 1 CPU node to drive the GPU
8 #SBATCH --partition=gpu # using the GPU on Midway3
9 #SBATCH --gres=gpu:1 # requesting only one GPU
10 #SBATCH --account=macs30123
11
12 module load cuda python
13
14 mpirun python3 ./<YOUR_FILE_NAME>.py
15
```

Assignment

Question 1,2 Review

- Question 1 (a) explicitly asks you to use pyCC, not jit.
- To make the code running, it would be helpful to have the below code.

```
1 if __name__ == "__main__":  
2     # YOUR CODE HERE
```

- Question 1 (b) asks you to compare the time speed comparison of up to using 20 cores, not more and not less.
- You should call the maximum number of nodes you are going to use in the sbatch.

Question 3 Objective

- **Objective:** compare the computation time and result between CPU and GPU
- Question (a): compare the computation time and result between CPU and GPU.
- Question (c): compare the computation time when the image is increased 10 times and 20 times.

Question 3 Cautions

- The main framework code and expected image of Q3 is given in the assignment description. You only need to make it into a code that can be run in OpenCL.
- The folder for accessing the input data are located in `/project2/mac30123/landsat8/`.
- **Only read the files in this directory, do not write on it.**
- **Restrict your CPU memory usage by 30G:** use `#SBATCH --mem-per-cpu=30G`
- **They are all in the assignment pdf, so please read carefully for the instructions before you run the code**

References i



Lisandro Dalcin. "MPI for Python". Accessed April 5, 2024.

<https://mpi4py.readthedocs.io/en/stable/index.html>



Khronos OpenCL Working Group. "The OpenCL Specification." Accessed April 5, 2024.

https://registry.khronos.org/OpenCL/specs/3.0-unified/html/OpenCL_API.html



PyOpenCL. "pyopencl 2024.1 documentation." Accessed April 5, 2024. <https://document.tician.de/pyopencl/>



Research Computing Center. "RCC User Guide." Accessed March 20, 2024. <https://rcc-uchicago.github.io/user-guide/>



SCHEDMD. "Slurm Workload Manager - sbatch." Accessed April 5, 2024. <https://slurm.schedmd.com/sbatch.html/>